

Subject: Regression Notes

Date: Friday, March 16, 2018 at 4:08:56 PM Pacific Daylight Time

From: Amar Pendala

To: Amar Pendala

Attachments: image001.png, image002.png, image003.png, image004.png, image005.png, image006.png, image007.png, image008.png, image009.png, image010.png, image011.png, image012.png, image013.png, image014.png, image015.png, image016.png, image017.png, image018.png, image019.png, image020.png, image021.png, image022.png, image023.png, image024.png, image025.png, image026.png, image027.png, image028.png, image029.png, image030.png, image031.png, image032.png, image033.png, image034.png, image035.png, image036.png, image037.png, image038.png

Week 1: Simple Regression

- Linear regression with a single feature
- Find a simple line that best fits the data

$$Y_i = W_0 + W_1 \cdot X_i + E_i$$

$$F(x) = W_0 + W_1 \cdot x$$

Week 2: Multiple Regression

- Linear regression w/ Multiple Features (ie. sq.ft, # of bathrooms, garage == house price)
- Most widely use tool out there in ML
- Polynomial Regression
 - Quadratic function: $F(x) = W_0 + W_1 \cdot X + W_2 \cdot X^2$
 - Higher Order function: $F(x) = W_0 + W_1 \cdot X + W_2 \cdot X^2 + \dots + W_N \cdot X^N$
 - $Y_i = W_0 + W_1 X_i + W_2 X_i^2 + \dots + W_N X_i^N + E_i$
 - Treat each power of X as a different **feature**
 - Feature 1 = 1 (constant) , parameter 1 = W_0
 - Feature 2 = X , parameter 2 = W_1
 - Feature 3 = X^2 , parameter 3 = W_2
 - Feature $N+1 = X^N$, parameter $N+1 = W_N$
 - Very useful in “detrending time series” (ie. values over time – house prices on average over time)
 - Also good to model seasonality (using sin/cosine features)
 - Housing sells best in summer
 - Weather modeling – varies annually & daily
 - Flu monitoring
 - Demand forecasting (jacket purchases)
- Notation (x_1 = sq ft, x_2 = # of bathrooms)
 - Additional notes:
 - \mathbf{x} = a vector that contains a list of features (ie. sq ft, bathrooms, etc...)
 - $\mathbf{x}[j]$ = get the jth item in the vector (ie. sq ft)
 - $\mathbf{x}[i]$ = vector for ith data point (ie. house i) in the data set
 - $\mathbf{x}_i[j]$ = get the jth item (ie. sq ft) in the vector for the ith data point (ie. house i)

General notation

Output: y ← scalar
 Inputs: $\mathbf{x} = (x[1], x[2], \dots, x[d])$
 ← d-dim vector

Notational conventions:

$x[j]$ = j^{th} input (scalar)
 $h_j(\mathbf{x})$ = j^{th} feature (scalar)
 \mathbf{x}_i = input of i^{th} data point (vector)
 $x_i[j]$ = j^{th} input of i^{th} data point (scalar)

So, simple hyperplane (ie. a curve without variations – ie weather model where it varies time of year, but also time of day)

$$y_i = w_0 + w_1 x_i[1] + \dots + w_d x_i[d] + \epsilon_i$$

feature 1 = 1

feature 2 = $x[1]$... e.g., sq. ft.

feature 3 = $x[2]$... e.g., #bath

...

feature $d+1$ = $x[d]$... e.g., lot size

For a d-dimensional curve (ie. a curve with variations – ie. weather model where it varies time of year, but also time of day)

- rather than a scalar to define the feature, you'd define a function (ie. sin/cosine) of $h_1(x_1)$

Model:

$$y_i = w_0 h_0(x_i) + w_1 h_1(x_i) + \dots + w_D h_D(x_i) + \epsilon_i$$

$$= \sum_{j=0}^D w_j h_j(x_i) + \epsilon_i$$

feature 1 = $h_0(\mathbf{x})$... e.g., 1

feature 2 = $h_1(\mathbf{x})$... e.g., $x[1]$ = sq. ft.

feature 3 = $h_2(\mathbf{x})$... e.g., $x[2]$ = #bath
 or, $\log(x[7])$ $x[2]$ = $\log(\text{\#bed}) \times \text{\#bath}$

...

- feature $D+1$ = $h_D(\mathbf{x})$... some other function of $x[1], \dots, x[d]$

observations $(\mathbf{x}_i, y_i) : N$

inputs $x[j] : d$

- # features $h_j(\mathbf{x}) : D$

■ Matrices & Vectors - addition/subtraction & multiplication:

- http://tutorial.math.lamar.edu/Classes/DE/LA_Matrix.aspx

- Special Matrices:

- Zero Matrix:

$$\mathbf{0}_{n \times m} = \begin{pmatrix} 0 & 0 & \dots & 0 \\ \vdots & \vdots & & \vdots \\ 0 & 0 & \dots & 0 \end{pmatrix}_{n \times m}$$

- identity Matrix:

$$\mathbf{I}_n = \begin{pmatrix} 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{pmatrix}_{n \times n}$$

- Adding matrices: (matrices must be same dimensions and add)

Example 1 Given the following two matrices,

$$A = \begin{pmatrix} 3 & -2 \\ -9 & 1 \end{pmatrix} \quad B = \begin{pmatrix} -4 & 1 \\ 0 & -5 \end{pmatrix}$$

compute $A-5B$.

Solution

There isn't much to do here other than the work.

$$\begin{aligned} A-5B &= \begin{pmatrix} 3 & -2 \\ -9 & 1 \end{pmatrix} - 5 \begin{pmatrix} -4 & 1 \\ 0 & -5 \end{pmatrix} \\ &= \begin{pmatrix} 3 & -2 \\ -9 & 1 \end{pmatrix} - \begin{pmatrix} -20 & 5 \\ 0 & -25 \end{pmatrix} \\ &= \begin{pmatrix} 23 & -7 \\ -9 & 26 \end{pmatrix} \end{aligned}$$

■

- Multiplying matrices (multiple row of Matrix A by column of Matrix B and add together)

Example 2 Given

$$A = \begin{pmatrix} 2 & -1 & 0 \\ -3 & 6 & 1 \end{pmatrix}_{2 \times 3} \quad B = \begin{pmatrix} 1 & 0 & -1 & 2 \\ -4 & 3 & 1 & 0 \\ 0 & 3 & 0 & -2 \end{pmatrix}_{3 \times 4}$$

compute AB .

Solution

The new matrix will have size 2×4 . The entry in row 1 and column 1 of the corresponding entries from the row of A and the column of B and then add them.

$$c_{11} = (2)(1) + (-1)(-4) + (0)(0) = 6$$

$$c_{12} = (2)(0) + (-1)(1) + (0)(3) = -1$$

$$c_{13} = (2)(-1) + (-1)(0) + (0)(-2) = -2$$

$$c_{14} = (2)(2) + (-1)(0) + (0)(-2) = 4$$

Here's the complete solution.

$$C = \begin{pmatrix} 6 & -1 & -2 & 4 \\ -27 & 21 & 9 & -8 \end{pmatrix}$$

■

- Determinant of matrices

$$\det(A) = |A|$$

re the formulas for the determinant of 2×2 and 3×3 matrices

$$\begin{vmatrix} a & c \\ b & d \end{vmatrix} = ad - cb$$

$$\begin{vmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{vmatrix} = a_{11} \begin{vmatrix} a_{22} & a_{23} \\ a_{32} & a_{33} \end{vmatrix} - a_{12} \begin{vmatrix} a_{21} & a_{23} \\ a_{31} & a_{33} \end{vmatrix} + a_{13} \begin{vmatrix} a_{21} & a_{22} \\ a_{31} & a_{32} \end{vmatrix}$$

■

OR

■

Now, notice that there are three diagonals that run from left to right and three diagonals that run from right to left. What we do is multiply the entries on each diagonal and then if the diagonal runs from left to right we add them up and if the diagonal runs from right to left we subtract them.

Here is the work for this matrix.

$$\begin{aligned} \det(B) &= \begin{vmatrix} 2 & 3 & 1 \\ -1 & -6 & 7 \\ 4 & 5 & -1 \end{vmatrix} = 2 \begin{vmatrix} -6 & 7 \\ 5 & -1 \end{vmatrix} - 3 \begin{vmatrix} -1 & 7 \\ 4 & -1 \end{vmatrix} + 1 \begin{vmatrix} -1 & -6 \\ 4 & 5 \end{vmatrix} \\ &= (2)(-6)(-1) + (3)(7)(4) + (1)(-1)(5) - \\ &\quad (3)(-1)(-1) - (2)(7)(5) - (1)(-6)(4) \\ &= 42 \end{aligned}$$

- Inverse Matrix

- If Matrix $AB = BA = I$ (Identity Matrix), then B is inverse of A ($B = A^{-1}$)
- To determine inverse of A , do this – create a new matrix of

■ Multiple Regression in Matrix format

Rewrite in matrix notation

For observation i

$$y_i = \sum_{j=0}^D w_j h_j(x_i) + \epsilon_i$$

$$\mathbf{y} = \begin{bmatrix} y_0 & y_1 & y_2 & \dots & y_D \end{bmatrix} = \begin{bmatrix} h_0(x_0) & h_1(x_0) & h_2(x_0) & \dots & h_D(x_0) \\ h_0(x_1) & h_1(x_1) & h_2(x_1) & \dots & h_D(x_1) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ h_0(x_D) & h_1(x_D) & h_2(x_D) & \dots & h_D(x_D) \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ \vdots \\ w_D \end{bmatrix} + \begin{bmatrix} \epsilon_0 \\ \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_D \end{bmatrix}$$

$$= \mathbf{w}^T \mathbf{h}(x_i) + \epsilon_i$$

○

For all observations together

$$\mathbf{y} = \mathbf{H} \mathbf{w} + \boldsymbol{\epsilon}$$

○

where \mathbf{H} is a matrix of transposed $\mathbf{h}(x)$ vectors

■ Residual sum of squares (RSS)

○ For simple linear regression:

$$\text{RSS}(w_0, w_1) = \sum_{i=1}^N (y_i - \underbrace{[w_0 + w_1 x_i]}_{\hat{y}_i(w_0, w_1)})^2$$

■

○ For multiple regression:

$$\text{RSS}(\mathbf{w}) = \sum_{i=1}^N (y_i - \underbrace{\mathbf{h}^T(x_i) \mathbf{w}}_{\hat{y}_i(\mathbf{w})})^2$$

2)

■

○ For multiple regression (in matrix notation):

$$\begin{aligned} \text{RSS}(\mathbf{w}) &= \sum_{i=1}^N (y_i - \mathbf{h}(x_i)^T \mathbf{w})^2 \\ &= (\mathbf{y} - \mathbf{H}\mathbf{w})^T (\mathbf{y} - \mathbf{H}\mathbf{w}) \end{aligned}$$

■

■ Gradient of RSS:

$$\begin{aligned} \nabla \text{RSS}(\mathbf{w}) &= \nabla [(\mathbf{y} - \mathbf{H}\mathbf{w})^T (\mathbf{y} - \mathbf{H}\mathbf{w})] \\ &= -2\mathbf{H}^T (\mathbf{y} - \mathbf{H}\mathbf{w}) \end{aligned}$$

○

■ Least squares D-Dimensional Curve

○ Closed Form Approach – set gradient to 0 & solve for \mathbf{w}

$$\nabla \text{RSS}(\mathbf{w}) = -2\mathbf{H}^T(\mathbf{y} - \mathbf{H}\mathbf{w}) = 0$$

Solve for \mathbf{w} :

$$\begin{aligned}
 & -2\mathbf{H}^T\mathbf{y} + 2\mathbf{H}^T\mathbf{H}\hat{\mathbf{w}} = 0 \\
 & \mathbf{H}^T\mathbf{H}\hat{\mathbf{w}} = \mathbf{H}^T\mathbf{y} \\
 & (\mathbf{H}^T\mathbf{H})^{-1} \mathbf{H}^T\mathbf{H}\hat{\mathbf{w}} = (\mathbf{H}^T\mathbf{H})^{-1} \mathbf{H}^T\mathbf{y} \\
 & \hat{\mathbf{w}} = (\mathbf{H}^T\mathbf{H})^{-1} \mathbf{H}^T\mathbf{y}
 \end{aligned}$$

$\bullet A^{-1}A = I$
 $\bullet I\mathbf{v} = \mathbf{v}$

- Complexity is $O(\text{features}^3)$ – very computationally expensive

○ **Gradient Descent approach (simpler) – most widely used algorithm in ML**

```

init  $\mathbf{w}^{(1)} = 0$  (or randomly, or smartly),  $t = 1$ 
while  $\|\nabla \text{RSS}(\mathbf{w}^{(t)})\| > \epsilon$ 
    for  $j = 0, \dots, D$ 
         $\text{partial}[j] = -2 \sum_{i=1}^N h_j(\mathbf{x}_i)(y_i - \hat{y}_i(\mathbf{w}^{(t)}))$ 
         $\mathbf{w}_j^{(t+1)} \leftarrow \mathbf{w}_j^{(t)} - \eta \text{partial}[j]$ 
     $t \leftarrow t + 1$ 

```

ϵ ← tolerance
 $\sqrt{\text{partial}[0]^2 + \dots + \text{partial}[D]^2}$